

# Advanced Data Modelling Techniques

Chris Travers

April 17, 2018

# Relational Databases

- Inferential systems
- Based on function composition
- Functions and function composition are the basis of normalization

# Object-Relational Databases

This course is about

- Intelligent datatypes
- Derived information
- Can extend to derive complex things

# Joins as Function Composition

- Relations are Functions
- In relations, fields are functions of keys.
- Foreign keys define function compositions
- Therefore we can define relations from joins

# Exercise 1

First exercise.

# what does this do?

```
select id, lock_record(r) from my_table r where id = 10;
```

What rows have lock\_record() run against them?

# Functional Programming

- FP is about passing functions around
- FP can optimize by composing functions
- FP can optimize when to run functions

(Go through Create Function syntax)

# Volatility

Volatility levels:

- Immutable
- Stable
- Volatile

(Live demo of stable and immutable on random).



# Functions and Indexes

- Functions can be indexed
- Expressions can be indexed
- These are run on write and can then sometimes be skipped on read

## Exercise 2

Second Exercise. Here we play around and look at expression-based indexes

## Discussion 2

Does it ever make sense to have an immutable function that reads from a table?

# Table Methods

PostgreSQL provides “table method” syntactic sugar

1. `t.val` checks to see if an attribute exists
2. if not, checks to see if a `val(t)` function exists

(Live Demo)

# Best Practices

- Make table methods immutable
- Since you have the whole row, use it.
- For compatibility with Pg before 9.4 (iirc) use \$1
- Understand the limits of function composition in Postgres

## Exercise 3

Imagine search that spans fields! we do this.

## Discussion 3

What are the pros/cons of this approach?

# Semistructured Data

- Often data is not perfectly structured
- Data may be present or not
- Structured text (EMBL or Uniprot in life sciences)
- Could be XML or JSON



# Use Cases

- Search on dates extracted from text files
- Search on emails as values in JSONB

(Live Demo)

## Exercise 4

Now let's do some text processing

# Discussion 4

Difficulties and dangers of semi-structured data

# Defining 1NF

- All relations have same number of fields
- All fields are atomic.

What does atomic mean? Discussion....

# Reasons to break 1NF

- Optimizing access patterns
- Ingestion
- Others?

# Design Options

- Arrays of Tuples
- JSON/JSONB
- Custom data types

(Life demos of all)

## Exercise 5

Now your turn to play with Non-1NF designs

Discussion: Problems and solutions of non-1NF designs.



# What is Inheritance in Pgsql

- Implicit casts
- Consistent interfaces

(Go through list of caveats regarding indexes)

# What is Inheritance Good For?

- Collective Management of Similar Tables
- Consistency in Data Types
- Common functional interfaces

(Live Demo)

## Exercise 6

Now to play with inheritance and notes.

# Multiple Inheritance

We can also use multiple inheritance for “reverse partitioning”  
(discuss use cases)

# Beware

This is not usual however, so:

- You may hit bugs (demonstrate type change bug)
- Column changes are hard
- Inheritance diamonds become actually useful

## Final Exercise

## Q/A on inheritance and multiple inheritance

Thanks everyone! Email me at [chris.travers@adjust.com](mailto:chris.travers@adjust.com) if you have any more questions or suggestions.  
(Further discussion available during hallway track).